

# Customer Churn Prediction

Customer Churn Prediction

```
[4]: #Importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classification_report
```

```
[8]: data=pd.read_csv('/Users/yashmantri/Downloads/Customer_data - customer_data.csv')
```

```
[10]: data.head()
```

```
[10]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection	TechSupp
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	No	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	No	
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...	Yes	

```
[12]: data.describe()
```

```
[12]:
```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7043.000000	7032.000000
mean	0.162147	32.371149	64.761692	2283.300441
std	0.368612	24.559481	30.090047	2266.771362
min	0.000000	0.000000	18.250000	18.800000
25%	0.000000	9.000000	35.500000	401.450000
50%	0.000000	29.000000	70.350000	1397.475000
75%	0.000000	55.000000	89.850000	3794.737500
max	1.000000	72.000000	118.750000	8684.800000

```
•[14]: #finding the number of null values from each column
data.isnull().sum()
```

```
[14]:
```

customerID	0
gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	11
Churn	0

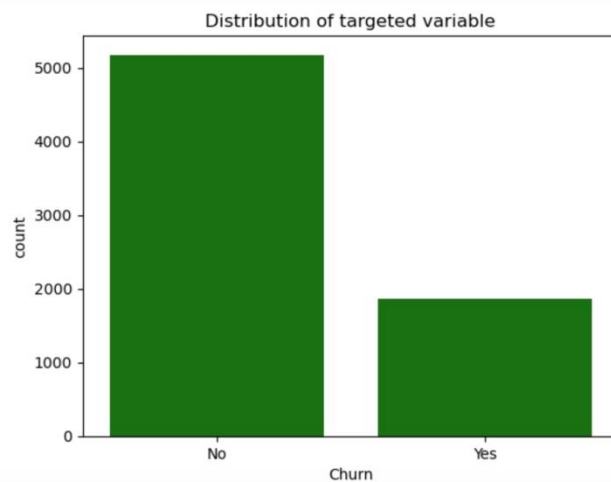
dtype: int64

```
[18]: #Replacinf Null Values with 0
data['TotalCharges']=data['TotalCharges'].fillna(0)
```

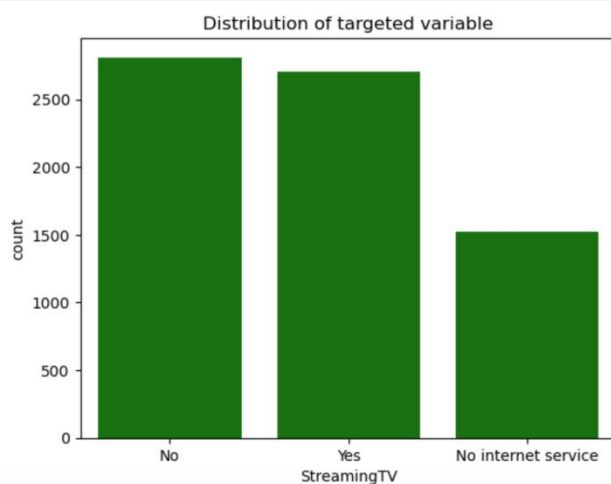
```
[21]: #fininf the number of null values from each column
data.isnull().sum()
```

```
[21]: customerID      0
gender            0
SeniorCitizen    0
Partner          0
Dependents       0
tenure           0
PhoneService     0
MultipleLines    0
InternetService  0
OnlineSecurity   0
OnlineBackup     0
DeviceProtection 0
TechSupport      0
StreamingTV      0
StreamingMovies  0
Contract         0
PaperlessBilling 0
PaymentMethod    0
MonthlyCharges   0
TotalCharges     0
Churn            0
dtype: int64
```

```
[27]: #Visulizing the distribution of targeted variable
sns.countplot(x='Churn',data=data, color='green')
plt.title('Distribution of targeted variable')
plt.show()
```



```
[39]: #Visulizing the distribution of targeted variable
sns.countplot(x='StreamingTV',data=data, color='green')
plt.title('Distribution of targeted variable')
plt.show()
```




```
[37]: #Correlation Matrix

num_list=list()
for columns in data.columns:
    if data[columns].dtype != object:
        num_list.append(columns)

correlation_matrix = data[num_list].corr()

#visulising the Matrix
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



```
[45]: bookmark=data.copy()
```

```
[53]: # Encode categorical variables
label_encoders = {}
for column in ['gender', 'Partner', 'Dependents', 'PhoneService',
               'MultipleLines', 'InternetService', 'OnlineSecurity',
               'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
               'PaperlessBilling', 'PaymentMethod', 'Churn']:
    le = LabelEncoder()
    data[column] = le.fit_transform(data[column])
    label_encoders[column] = le
```

```
[55]: data
```

```
[55]: OnlineSecurity ... DeviceProtection TechSupport StreamingTV StreamingMovies Contract PaperlessBilling PaymentMethod MonthlyCharges TotalCharges Churn
0 ... 0 0 0 0 0 0 1 2 29.85 29.85 0
2 ... 2 0 0 0 0 1 0 3 56.95 1889.50 0
2 ... 0 0 0 0 0 0 1 3 53.85 108.15 1
2 ... 2 2 0 0 1 0 0 0 42.30 1840.75 0
0 ... 0 0 0 0 0 0 1 2 70.70 151.65 1
... ... ... ... ... ... ... ... ... ... ...
2 ... 2 2 2 2 1 1 3 84.80 1990.50 0
0 ... 2 0 2 2 1 1 1 103.20 7362.90 0
2 ... 0 0 0 0 0 1 2 29.60 346.45 0
0 ... 0 0 0 0 0 1 3 74.40 306.60 1
2 ... 2 2 2 2 2 1 0 105.65 6844.50 0
```

```

[59]: # Features and target
X = data.drop(['customerID', 'Churn'], axis=1)
y = data['Churn']

[63]: # Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)

[67]: # Feature scaling
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

• [69]: # Logistic Regression

logreg = LogisticRegression()
logreg.fit(X_train, y_train)

logreg_pred = logreg.predict(X_test)

[71]: logreg_accuracy = accuracy_score(y_test, logreg_pred)
logreg_precision = precision_score(y_test, logreg_pred)
logreg_recall = recall_score(y_test, logreg_pred)
logreg_f1score = f1_score(y_test, logreg_pred)

[75]: # Display accuracy and performance metrics
print(f'Logistic Regression Accuracy: {(logreg_accuracy) * 100:.2f}%')
print(f'Logistic Regression Precision: {(logreg_precision) * 100:.2f}%')
print(f'Logistic Regression Recall: {(logreg_recall) * 100:.2f}%')
print(f'Logistic Regression F1score: {(logreg_f1score) * 100:.2f}%')

Logistic Regression Accuracy: 79.79%
Logistic Regression Precision: 63.35%
Logistic Regression Recall: 54.07%
Logistic Regression F1score: 58.34%

```